Before starting:
    - Confirm that we have the user's signed consent form
    - Confirm that they can see both notebooks

*** Interview Starts ***

R: Hi, thank you for participating in this user study. Let's start with some basic information. Can you please tell me:

* Your first name: _____
* Your age: _____
* Your gender: _____
* Your occupation: _____
* How many years of experience do you have working in data science/machine learning: _____

R: Now, I'm going to ask you a few questions to set a baseline for your familiarity with some concepts we'll be exploring in this study. I'm going to list a few concepts of statistcs and machine learning. For each of them please tell on a scale of 1 to 5, where 1 is not familiar at all and 5 is very familiar, how familiar you are with each of the following:

* Posterior probability
* Log Likelihood
* Log Odds
* Weight of Evidence Score
* Local interpretability methods like LIME or SHAP

*** Part 1 ***

* Q1: What aspect of the news article is the model relying on the most to make its prediction here?
* Q2: The model is not too confident (p=0.68) that this news article belongs to the 'entertainment' class. Why do you think that is? How would you change the article to cause the model to be more confident?

*** Part 2 ***

* Q3: Why do you think the model didn't predict 'business' instead?
* Q4: Why do you think the model didn't predict 'social media' instead?
* Q5: For other news articles that have similar keywords to this one, how much more likely do you think it is that the model with predict `world` instead of other classes?

*** Follow-Up Questions ***

* Could you rate your experience using the interpretability tool (1-5)
* Did you find the option to aggregate features intro groups useful? Yes/No, Why?
* Did you find the option to produce multi-step explanations useful? Yes/No, Why?